HW1

Gabriel ROMON

1 DCT denoiser

1. The conditional density of X given $Y = y$ is given by $f_{X|Y=y}(x) = \dfrac{f_{Y|X=x}(y)f_X(x)}{f_Y(y)}$.

   Since $Y = X + B$, $Y|X = x \sim \mathcal{N}(x, \sigma^2)$ hence $f_{X|Y=y}(x) = \frac{1}{\sqrt{2\pi\sigma^2}}\exp(-\frac{1}{2\sigma^2}(y-x)^2)$. Thus

   $$f_{X|Y=y}(x) \propto \exp(-\frac{1}{2\sigma^2}(y-x)^2 - |x|)$$

   Computing the maximum of $f_{X|Y=y}(x)$ is consequently equivalent to computing

   $$\arg\max_{x\in\mathbb{R}} \exp\left(-\frac{1}{2\sigma^2}(y-x)^2 - |x|\right) = \arg\min_{x\in\mathbb{R}} \frac{1}{2\sigma^2}(y-x)^2 + |x|$$

   Let $\varphi : x \mapsto \frac{1}{2\sigma^2}(y-x)^2 + |x|$. Studying the derivative of φ on $\mathbb{R}^+$ and $\mathbb{R}^-$ shows that if $y \geq \sigma^2$, φ is increasing over $[y-\sigma^2, \infty)$ and decreasing on the complement. If $y \leq -\sigma^2$, φ is decreasing over $(-\infty, y+\sigma^2]$ and increasing on the complement. Therefore, if $y \geq \sigma^2$, the minimum is attained at $y-\sigma^2$, if $y \in (-\sigma^2, \sigma^2)$, the minimum is attained at 0 and if $y \leq -\sigma^2$, the minimum is attained at $y+\sigma^2$. This rewrites more compactly as

   $$x^* = \left(1 - \frac{\sigma^2}{|y|}\right)^+ y$$

   which corresponds to *soft-thresholding* of y with threshold σ^2.

2. The code in `DCT_denoiser` performs denoising on an image corrupted by Gaussian noise with variance σ^2.

   First, a tensor `D` of shape $N \times N \times N^2$ is built such that the `D[:,:,k]` are the inverse discrete cosine transforms of the N^2 elementary matrices of $\mathbb{R}^{N\times N}$. More precisely there is some bijection $\varphi : [\![1, N]\!]^2 \to [\![1, N^2]\!]$ such that $D[:,:,\varphi(i,j)]$ is the 2D inverse DCT of the elementary matrix E_{ij}. Consequently, $D[m,n,\varphi(i,j)] = \alpha_i\alpha_j \cos\left(\frac{\pi i}{N}\left(m+\frac{1}{2}\right)\right)\cos\left(\frac{\pi j}{N}\left(n+\frac{1}{2}\right)\right)$ where the α_i are the usual constants defining the type 2 orthonormal DCT.

   For some fixed (i,j), $D[:,:,\varphi(i,j)]$ is a kernel/filter. For each $N \times N$ patch X of the noisy image, `conv2D` computes the correlation between the filter $D[:,:,\varphi(i,j)]$ and X:

   $$\sum_{m=0}^{N-1}\sum_{n=0}^{N-1} D[m,n,\varphi(i,j)]X_{mn} = \alpha_i\alpha_j \sum_{m=0}^{N-1}\sum_{n=0}^{N-1} X_{mn}\cos\left(\frac{\pi i}{N}\left(m+\frac{1}{2}\right)\right)\cos\left(\frac{\pi j}{N}\left(n+\frac{1}{2}\right)\right)$$

   which is precisely coefficient (i,j) in the 2D DCT of the patch X.

   This is done for every (i,j) and the resulting coefficients undergo *hard-thresholding*: all coefficients with absolute value larger than 3σ are set to 0. This is different from the soft-thresholding used in 1. Next, the algorithm returns to the spatial domain by convolving the resulting tensor with the same filters (note that this an actual convolution so filters have to flipped beforehand, hence the command `np.fliplr(np.flipud(D))`).

3. In order to learn the threshold parameter s, one has to define a loss function and get some training data. Regarding the loss, a possible choice is the squared euclidean norm on the flattened images $\ell(y, \hat{y}) = \sum_i (y_i - \hat{y}_i)^2$. However, the relation between ℓ and s is hard to write down explicitly (even as a composition of several simple functions) thus making gradient descent difficult.

4. Computing the inverse DCT of the elementary matrices has total cost $O(N^2 \times N\log N) = O(N^3\log N)$. For an $n \times m$ input, the correlation operation has cost $O(N^2 \times (m-N)(n-N) \times N^2) = O(nmN^4)$ since for each filter, there are $\sim (m-N)(n-N)$ patches and $O(N^2)$ operations to compute on each. The second convolution has the same cost $O(nmN^4)$, so the number of operations per pixel is $O(N^4)$.